


Classical and Quantum Theory for QPU

From Bits to Bloch Spheres



Bailey Forbes
September 13, 2026

A ground-up explanation of classical logic, complex amplitudes, quantum state, measurement, entanglement, and the mathematics represented by QPU gates.

Contents

1	Classical computation foundations	1
1.1	Classical bits and logical state	1
1.2	Boolean algebra	1
1.3	Truth tables as complete finite specifications	1
1.4	Combinational logic and stored state	2
1.5	Why ordinary classical gates can lose information	2
1.6	Controls, targets, and workspace	2
1.7	Fan-out, copying, and the no-cloning boundary	2
1.8	Deterministic, probabilistic, and quantum models	3
1.9	Classical and QPU operation map	3
2	Mathematical foundations	3
2.1	Real numbers and complex numbers	3
2.2	The complex plane, magnitude, and conjugate	4
2.3	Angles, radians, sine, and cosine	4
2.4	Euler’s formula and Euler’s identity	4
2.5	Vectors and basis vectors	5
2.6	Normalization and the Born rule	5
2.7	Matrices as gates	5
2.8	Unitarity and reversibility	5
3	Qubits, phase, and the Bloch sphere	6
3.1	A qubit is not just an unknown bit	6
3.2	Global phase and relative phase	6
3.3	Bloch-sphere coordinates	6
3.4	How gates move a Bloch vector	6
4	Multiple qubits and entanglement	7
4.1	Tensor products	7
4.2	Entanglement	7
4.3	Controlled operations	7
5	Measurement, probability, and truth tables	8
5.1	Projective measurement	8
5.2	Single runs and repeated runs	8
5.3	Truth tables versus quantum amplitudes	8
5.4	A complete interference example	8
6	Single-qubit primitive gates	9
6.1	Pauli X ($\mathbf{X}$)	9
6.2	Pauli Y ($\mathbf{Y}$)	9
6.3	Pauli Z ($\mathbf{Z}$)	9

6.4	Hadamard (H)	10
6.5	S phase (S)	10
6.6	T phase (T)	11
6.7	Arbitrary phase (PHASE)	11
6.8	Dagger and inverse syntax	11
7	Controlled and two-qubit primitive gates	12
7.1	Controlled NOT (CNOT)	12
7.2	Doubly controlled NOT / Toffoli (CCNOT)	12
7.3	Controlled Z (CZ)	13
7.4	Controlled Y (CY)	13
7.5	Swap (SWAP)	14
8	Derived logical gates	14
8.1	Logical NOT (NOT)	14
8.2	Reversible AND (AND)	15
8.3	Reversible NAND (NAND)	15
8.4	Reversible OR (OR)	16
8.5	Reversible XOR (XOR)	16
9	Measurement gate	17
9.1	Measurement (MEASURE)	17
10	Reset: an operation that is not a gate	17
10.1	Reversible gates versus RESET	17
10.2	Where RESET comes from	18
10.3	Reset buttons compared	18
11	Glossary	18

1 Classical computation foundations

Quantum circuits build on ideas from classical digital logic. Understanding bits, Boolean functions, truth tables, and reversible computation makes the QPU gate model much easier to use.

1.1 Classical bits and logical state

A classical bit has one definite logical value:

$$b \in \{0, 1\}.$$

Physical computers may represent that value with voltage, charge, magnetic orientation, or another measurable quantity. Digital logic ignores most physical detail and works with the abstraction 0 or 1.

A wire carries a bit from one operation to another. A *gate* computes a function of one or more incoming bits. For example, NOT has one input and one output:

$$\text{NOT}(0) = 1, \quad \text{NOT}(1) = 0.$$

In QPU source, a token names the logical wire. On computational-basis states, $X \text{ } -I \text{ } Q \text{ } -0 \text{ } Q$ has the same measured input/output relationship as classical NOT.

1.2 Boolean algebra

Boolean algebra describes functions whose values are 0 and 1. This guide uses:

Operation	Symbol	Spoken form	Meaning
NOT	$\neg A$	not A	1 when A is 0; 0 when A is 1.
AND	$A \wedge B$	A and B	1 only when both inputs are 1.
OR	$A \vee B$	A or B	1 when either or both inputs are 1.
XOR	$A \oplus B$	A exclusive-or B	1 when exactly one input is 1.
NAND	$\neg(A \wedge B)$	not-and	The complement of AND.

For bits, XOR is addition modulo two:

$$0 \oplus 0 = 0, \quad 0 \oplus 1 = 1, \quad 1 \oplus 0 = 1, \quad 1 \oplus 1 = 0.$$

There is no carry in modulo-two addition. Ordinary binary addition handles the carry separately, which is why a half-adder produces both Sum and Carry.

1.3 Truth tables as complete finite specifications

A truth table lists a function's output for every possible input pattern. With n binary inputs there are 2^n patterns. AND, OR, XOR, and NAND can be compared in one table:

A	B	$A \wedge B$	$A \vee B$	$A \oplus B$	$\neg(A \wedge B)$
0	0	0	0	0	1
0	1	0	1	1	1
1	0	0	1	1	1
1	1	1	1	0	0

For a deterministic combinational circuit, this table completely specifies the classical input/output behavior. It does not reveal the circuit's internal gate arrangement, propagation time, power consumption, or quantum phase.

1.4 Combinational logic and stored state

A *combinational* output depends only on current inputs. Adders and the Boolean gates above are combinational. A *sequential* circuit also depends on stored state from an earlier step, as in counters and latches.

QPU cycle suffixes can document successive stages, and INCREASECYCLE marks a workspace boundary, but the compiled circuit is still one ordered gate sequence. `SAVE_STATE` and `LOAD_STATE` copy and restore the simulator state vector. They are not quantum gates. When modeling stateful behavior inside the unitary circuit, expose previous state as `PARAMS` and return next state through `RETURNVALS`.

1.5 Why ordinary classical gates can lose information

A function is reversible only if each output identifies exactly one input. Classical NOT is reversible: seeing output 1 proves the input was 0. AND is not: output 0 could have come from 00, 01, or 10.

$$(A, B) \xrightarrow{\text{AND}} A \wedge B$$

maps four input patterns onto only two output values, discarding information. Quantum evolution through closed-system gates must be unitary and therefore reversible. A quantum/reversible AND preserves both inputs and mutates another target:

$$(A, B, t) \mapsto (A, B, t \oplus (A \wedge B)).$$

Knowing the output triple lets the transformation be run backward to recover the input triple.

How to read $t' = t \oplus (A \wedge B)$

- A and B are **controls**. They are only read. After the gate they still hold exactly what they held before, which is why the input can always be recovered.
- t is the **target**, the output workspace. It is the only wire that changes. The prime in t' means “ t after the gate”.
- $\oplus$ means “flip if”. The target flips when $A \wedge B$ is 1 and stays put otherwise.
- Start the target at 0 and it ends holding $A \wedge B$, the ordinary AND answer. Start it at 1 and it ends holding NAND instead.

In the Circuit builder, A is **Control A**, B is **Control B**, and t is **Target particle**.

1.6 Controls, targets, and workspace

These terms connect classical logic to QPU syntax:

Control A wire whose value decides whether a target operation occurs. In `CNOT -I A -O T`, A is the control.

Target The wire changed by an operation. T is the target in that `CNOT`.

Workspace or ancilla An extra wire introduced to hold an intermediate or output while preserving inputs. It is commonly initialized to $|0\rangle$.

Garbage Intermediate information left in workspace after the desired answer has been computed. Reversible algorithms often *uncompute* temporary gates to return garbage wires to zero.

Uncomputation Applying inverse operations in reverse order to clean workspace without erasing information irreversibly.

The QPU language calls logical wire names tokens. `CREATETOKEN` reserves workspace, `SET` prepares it, `-I` lists inputs/controls, and `-O` names targets.

1.7 Fan-out, copying, and the no-cloning boundary

Classical logic freely copies a bit onto several wires. If the source is a known basis bit and targets begin at zero, `CNOT` reproduces that basis value:

$$|x\rangle|0\rangle \xrightarrow{\text{CNOT}} |x\rangle|x\rangle, \quad x \in \{0, 1\}.$$

For a superposition $\alpha|0\rangle + \beta|1\rangle$, the same circuit produces

$$\alpha|00\rangle + \beta|11\rangle,$$

not

$$(\alpha|0\rangle + \beta|1\rangle) \otimes (\alpha|0\rangle + \beta|1\rangle).$$

The first state is generally entangled; the second would be two independent copies and contains cross terms. The no-cloning theorem forbids a universal gate that copies every unknown quantum state. QPU aliases similarly give two names to one wire; they do not create a physical copy.

1.8 Deterministic, probabilistic, and quantum models

Model	State description	Transformation	Observation
Deterministic classical	One definite bit string	Boolean function	The definite output bits
Probabilistic classical	Probabilities over bit strings	Stochastic update	A sampled string
Pure-state quantum	Complex amplitudes over basis strings	Unitary matrix	A sampled string after magnitude-squared probabilities and collapse

Negative and imaginary amplitudes are not negative or imaginary probabilities. They carry phase. Two computational paths can add constructively or cancel destructively before probabilities are calculated. This interference has no counterpart in an ordinary classical random-choice table.

1.9 Classical and QPU operation map

Classical idea	QPU operation	Important difference
NOT	X or NOT	Acts linearly on superpositions; X is unitary.
XOR into target	CNOT	Preserves the control and can entangle.
AND into target	CCNOT or AND	Preserves controls; target is XOR-mutated.
Swap variables	SWAP	Exchanges complete quantum wire states.
Random bit	H then MEASURE	H first creates coherent amplitudes, not an ordinary hidden random choice.
Read a bit	MEASURE	Sampling collapses a superposition and can disturb entanglement.
Initialize zero	SET token 0p	Compiles to state preparation/reset behavior.
Function output	RETURNVALS	Declares an interface; does not itself copy or measure a wire.

2 Mathematical foundations

This chapter assumes no prior quantum mechanics. The essential ideas are numbers, vectors, probability, and transformations. Quantum notation can look unfamiliar, but each symbol has a concrete computational meaning.

2.1 Real numbers and complex numbers

The real numbers include familiar values such as -2 , 0 , $1/3$, and $\sqrt{2}$. A real number can be placed on a one-dimensional number line. Quantum amplitudes require a second numerical direction, represented by the *imaginary unit*

$$i = \sqrt{-1}, \quad i^2 = -1.$$

There is no real number whose square is -1 , so i extends the number system. A complex number has the form

$$z = a + bi,$$

where a is the real part and b is the imaginary part. Examples are $3 + 2i$, $-i$, and 1 (which is $1 + 0i$).

Complex arithmetic follows ordinary algebra together with $i^2 = -1$:

$$(a + bi) + (c + di) = (a + c) + (b + d)i,$$

$$(a + bi)(c + di) = (ac - bd) + (ad + bc)i.$$

For example,

$$(1 + i)^2 = 1 + 2i + i^2 = 2i.$$

2.2 The complex plane, magnitude, and conjugate

A complex number is a point or arrow in a plane: horizontal coordinate a , vertical coordinate b . Its magnitude is the arrow length,

$$|z| = \sqrt{a^2 + b^2}.$$

The complex conjugate reflects the arrow across the real axis:

$$z^* = a - bi.$$

Multiplying by the conjugate gives a nonnegative real number:

$$z^*z = (a - bi)(a + bi) = a^2 + b^2 = |z|^2.$$

Quantum measurement probabilities use $|z|^2$, not z itself. This guarantees real, nonnegative probabilities even though amplitudes may be complex.

2.3 Angles, radians, sine, and cosine

An arrow in the complex plane may be described by its length r and angle θ :

$$a = r \cos \theta, \quad b = r \sin \theta.$$

Quantum formulas normally measure angles in radians. One full turn is 2π radians, a half turn is π , and a quarter turn is $\pi/2$. Conversion is

$$\theta_{\text{radians}} = \theta_{\text{degrees}} \frac{\pi}{180}.$$

This is why PHASE=90d and PHASE=pi/2 describe the same angle.

2.4 Euler's formula and Euler's identity

Euler's formula connects complex numbers to rotations:

$$e^{i\theta} = \cos \theta + i \sin \theta.$$

The value $e^{i\theta}$ always has magnitude one. As θ changes, it moves around the unit circle in the complex plane. Important cases are

$$e^{i0} = 1, \quad e^{i\pi/2} = i, \quad e^{i\pi} = -1, \quad e^{i2\pi} = 1.$$

Setting $\theta = \pi$ gives Euler's identity,

$$e^{i\pi} + 1 = 0.$$

In a phase gate, multiplying an amplitude by $e^{i\theta}$ rotates that amplitude without changing its magnitude. Its immediate measurement probability is unchanged, but its relationship to other amplitudes changes. Later gates can convert that relative phase into different probabilities.

2.5 Vectors and basis vectors

A column vector is an ordered list of numbers. The two basis vectors used for one qubit are

$$|0\rangle = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \quad |1\rangle = \begin{bmatrix} 0 \\ 1 \end{bmatrix}.$$

The vertical bars and angle bracket form Dirac's *ket* notation. `0p` prepares $|0\rangle$ and `1p` prepares $|1\rangle$. Any one-qubit pure state is a complex linear combination:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle = \begin{bmatrix} \alpha \\ \beta \end{bmatrix}.$$

The complex numbers α and β are *probability amplitudes*. They are not probabilities themselves.

2.6 Normalization and the Born rule

A valid state is normalized:

$$|\alpha|^2 + |\beta|^2 = 1.$$

The Born rule turns amplitudes into measurement probabilities:

$$P(0) = |\alpha|^2, \quad P(1) = |\beta|^2.$$

For the plus state,

$$|+\rangle = \frac{|0\rangle + |1\rangle}{\sqrt{2}},$$

both amplitudes are $1/\sqrt{2}$. Squaring their magnitudes gives $P(0) = P(1) = 1/2$. The factor $1/\sqrt{2}$ is therefore required; writing $|0\rangle + |1\rangle$ without normalization would produce total probability two.

2.7 Matrices as gates

A matrix is a rectangular array of numbers. A one-qubit gate is represented by a 2×2 complex matrix. Matrix-vector multiplication computes the new amplitudes:

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} \alpha \\ \beta \end{bmatrix} = \begin{bmatrix} a\alpha + b\beta \\ c\alpha + d\beta \end{bmatrix}.$$

For example, the X gate exchanges the two amplitudes:

$$\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} \alpha \\ \beta \end{bmatrix} = \begin{bmatrix} \beta \\ \alpha \end{bmatrix}.$$

Applying gates in source order means multiplying by matrices in the same temporal order: if U acts first and V second, then

$$|\psi_{\text{final}}\rangle = V(U|\psi_{\text{initial}}\rangle) = VU|\psi_{\text{initial}}\rangle.$$

2.8 Unitarity and reversibility

Closed-system quantum gates must preserve total probability. Their matrices are *unitary*:

$$U^\dagger U = I.$$

Here $U^\dagger$ means transpose the matrix and complex-conjugate every entry; I is the identity matrix. Unitarity guarantees an inverse, $U^{-1} = U^\dagger$. Quantum gates therefore cannot normally erase information. Classical functions such as AND are embedded reversibly by preserving inputs and XORing the result into a target:

$$|a, b, t\rangle \mapsto |a, b, t \oplus (a \wedge b)\rangle.$$

Measurement and reset are exceptions because they interact with classical outcomes or preparation rather than acting as closed unitary gates.

3 Qubits, phase, and the Bloch sphere

3.1 A qubit is not just an unknown bit

A classical bit is exactly 0 or 1. A qubit may be in any normalized superposition $\alpha|0\rangle + \beta|1\rangle$. This does not mean that the qubit secretly chose a classical value before measurement. Relative amplitudes can interfere, producing effects no ordinary random bit can reproduce.

3.2 Global phase and relative phase

Multiplying an entire state by a unit complex number does not change physical predictions:

$$|\psi\rangle \quad \text{and} \quad e^{i\gamma}|\psi\rangle$$

have the same measurement probabilities and are considered equivalent up to a *global phase*. Relative phase is different. The states

$$|+\rangle = \frac{|0\rangle + |1\rangle}{\sqrt{2}}, \quad |-\rangle = \frac{|0\rangle - |1\rangle}{\sqrt{2}}$$

give equal 0/1 probabilities if measured immediately, yet H sends $|+\rangle \mapsto |0\rangle$ and $|-\rangle \mapsto |1\rangle$. Their relative sign is observable through interference.


3.3 Bloch-sphere coordinates

After removing an unobservable global phase, every pure one-qubit state can be written with two real angles:

$$|\psi\rangle = \cos\left(\frac{\theta}{2}\right)|0\rangle + e^{i\phi}\sin\left(\frac{\theta}{2}\right)|1\rangle,$$

where $0 \leq \theta \leq \pi$ and $0 \leq \phi < 2\pi$. These angles define a point on the unit Bloch sphere:

$$x = \sin\theta \cos\phi, \quad y = \sin\theta \sin\phi, \quad z = \cos\theta.$$



The north pole is $|0\rangle$ and the south pole is $|1\rangle$. The equator contains equal-magnitude superpositions. $|+\rangle$ points along $+x$, $|-\rangle = (|0\rangle - |1\rangle)/\sqrt{2}$ along $-x$, and $|+i\rangle = (|0\rangle + i|1\rangle)/\sqrt{2}$ along $+y$.

3.4 How gates move a Bloch vector

Ignoring global phase, many one-qubit gates are rotations:

- X rotates the sphere by π around the x-axis, exchanging north and south.

- Y rotates by π around the y-axis.
- Z rotates by π around the z-axis, changing relative phase.
- S and T rotate around z by $\pi/2$ and $\pi/4$.
- PHASE rotates around z by its supplied angle θ .
- H rotates the coordinate description so the z basis and x basis are exchanged: $|0\rangle \leftrightarrow |+\rangle$ and $|1\rangle \leftrightarrow |-\rangle$.

The sphere is a teaching model for one pure qubit. It cannot display the full state of entangled multi-qubit systems as one ordinary three-dimensional arrow.

4 Multiple qubits and entanglement

4.1 Tensor products

Two classical bits have four possibilities; two qubits have four basis vectors:

$$|00\rangle, |01\rangle, |10\rangle, |11\rangle.$$

Independent systems combine with the tensor product $\otimes$. If $|A\rangle = [a_0, a_1]^T$ and $|B\rangle = [b_0, b_1]^T$, then

$$|A\rangle \otimes |B\rangle = \begin{bmatrix} a_0 b_0 \\ a_0 b_1 \\ a_1 b_0 \\ a_1 b_1 \end{bmatrix}.$$

For n qubits, the simulator stores 2^n amplitudes:

$$|\Psi\rangle = \sum_{x=0}^{2^n-1} a_x |x\rangle, \quad \sum_x |a_x|^2 = 1.$$

This exponential growth is why state-vector simulation becomes expensive as qubit count increases.

4.2 Entanglement

Some multi-qubit states cannot be factored into independent one-qubit states. The Bell state

$$|\Phi^+\rangle = \frac{|00\rangle + |11\rangle}{\sqrt{2}}$$

is entangled. Each qubit measured alone is equally likely to be 0 or 1, but their outcomes are perfectly correlated. It can be prepared from $|00\rangle$:

```
H -I A -0 A
CNOT -I A -0 B
```

The first gate creates a superposition; the controlled gate correlates B with A. Entanglement is not a classical alias. It is a property of the joint amplitude vector.

Correlated is not the same as copied

After the *Bell state* starter circuit, both wires always measure the same value, 00 or 11, each half the time. That looks like copying, but it is not. Before measurement neither wire has a value of its own: stop just before the first meter with **Step gate**, and the **Final state vector** lists only basis states in which q0 and q1 agree, with no separate state for either wire. Choose **Reset state** and **Run all** several times and the pair changes together between 00 and 11.

4.3 Controlled operations

A controlled gate applies a target operation only on basis components where its control condition is true. CNOT applies X when its control is one. If the control is in superposition, linearity applies the operation to one branch and

not the other, which can create entanglement:

$$\frac{|0\rangle + |1\rangle}{\sqrt{2}} \otimes |0\rangle \xrightarrow{\text{CNOT}} \frac{|00\rangle + |11\rangle}{\sqrt{2}}.$$

CCNOT requires both controls to be one. The system’s OR and XOR derived gates use “any active” and “odd parity” predicates respectively.

5 Measurement, probability, and truth tables

5.1 Projective measurement

For a measured qubit q in a multi-qubit state,

$$P(q = 1) = \sum_{x: x_q=1} |a_x|^2, \quad P(q = 0) = 1 - P(q = 1).$$

The simulator samples one outcome. Amplitudes incompatible with that outcome become zero; compatible amplitudes are divided by the square root of the kept probability so the state is normalized again. Measurement is destructive in the sense that the pre-measurement superposition is generally unavailable after collapse.

5.2 Single runs and repeated runs

One simulation run produces one measurement sample. To estimate a probability, run the same preparation and circuit many times. If a plus state is measured 1 approximately 500 times in 1000 shots, the observed frequency 500/1000 estimates the theoretical probability 1/2. More shots reduce ordinary sampling noise but never reveal a complex amplitude directly.

5.3 Truth tables versus quantum amplitudes

A classical truth table specifies output bits for listed input bits. It is a good contract for arithmetic and reversible Boolean circuits operating on computational-basis states. It cannot fully describe phase, interference, entanglement, or a probability distribution.

The QPU truth-table runner prepares each input row, executes the circuit, measures declared RETURNVALS, and compares those measured bits. In companion `.qpuio` tables, cells are `0p`, `1p`, or `sp`. On an expected output, `sp` means “either measured bit is acceptable.” Use state equations and repeated experiments in addition to tables when a circuit’s purpose is genuinely quantum.

5.4 A complete interference example

Starting in zero, two Hadamards restore zero:

$$|0\rangle \xrightarrow{H} |+\rangle \xrightarrow{H} |0\rangle.$$

Inserting Z changes the relative phase and changes the final answer:

$$|0\rangle \xrightarrow{H} |+\rangle \xrightarrow{Z} |-\rangle \xrightarrow{H} |1\rangle.$$

```
MAIN-PROCESS Interference
CREATETOKEN -I Q
SET Q 0p
H -I Q -0 Q
Z -I Q -0 Q
H -I Q -0 Q
MEASURE -I Q
RETURNVALS Q
```


The middle Z does not change an immediate 0/1 probability, yet interference at the final H makes its effect deterministic. This is the central reason the system documents matrices and phase as well as truth tables.

6 Single-qubit primitive gates

Placing a single-qubit gate

X, Y, Z, H, S, T, and PHASE act on one wire, so the workbench uses only **Target particle**; **Control A** and **Control B** are ignored. PHASE also shows a **Phase angle** slider. On the canvas, drop the block on the wire you want to change.

6.1 Pauli X (X)




$$X = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}, \quad X|0\rangle = |1\rangle, \quad X|1\rangle = |0\rangle, \quad X(\alpha|0\rangle + \beta|1\rangle) = \beta|0\rangle + \alpha|1\rangle.$$

X is the quantum bit flip. On the Bloch sphere it is a half-turn around the x-axis. It exchanges the north and south poles and is self-inverse: $X^2 = I$. Unlike a classical NOT implementation, X remains meaningful on a superposition because it exchanges amplitudes linearly without measuring.

```
X -I Q:0 -0 Q:0
```

Input	Classical comparison	Output
0	NOT 0	1
1	NOT 1	0

6.2 Pauli Y (Y)




$$Y = \begin{bmatrix} 0 & -i \\ i & 0 \end{bmatrix}, \quad Y|0\rangle = i|1\rangle, \quad Y|1\rangle = -i|0\rangle.$$

It flips the basis value like X but adds phases that become observable through later interference. Geometrically it is a half-turn around the y-axis and is self-inverse, $Y^2 = I$. The factors i and $-i$ are global phases for either basis state by itself, but on components of a larger superposition they may become relative phases.

```
Y -I Q -0 Q
```

Input	Measured bit	Quantum output
0	1	$i 1\rangle$
1	0	$-i 0\rangle$

6.3 Pauli Z (Z)




$$Z = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}, \quad Z|0\rangle = |0\rangle, \quad Z|1\rangle = -|1\rangle.$$

Z does not change a computational-basis measurement, but it changes relative phase. For example, $Z|+\rangle = |-\rangle = (|0\rangle - |1\rangle)/\sqrt{2}$. On the Bloch sphere it is a half-turn around z. It is self-inverse and is the simplest example showing why quantum state cannot be reduced to a table of immediate measurement probabilities.

```
Z -I Q -0 Q
```

Input	Classical bit	Quantum output
0	0	$ 0\rangle$
1	1	$- 1\rangle$

6.4 Hadamard (H)



$$H = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}, \quad H|0\rangle = |+\rangle, \quad H|1\rangle = |-\rangle.$$


Hadamard changes between the computational z basis and the plus/minus x basis. It creates equal magnitudes from a basis state, but the sign of the second amplitude depends on the input. It is also self-inverse: applying H twice returns the original state. H is commonly used before controlled gates to create entanglement and after phase gates to reveal interference.

```
SET Q 0p
H -I Q -0 Q
MEASURE -I Q
```

Input	Measurement distribution	State before measurement
0	50% 0, 50% 1	$(0\rangle + 1\rangle)/\sqrt{2}$
1	50% 0, 50% 1	$(0\rangle - 1\rangle)/\sqrt{2}$

Equal classical probabilities hide the relative minus sign responsible for interference.

6.5 S phase (S)




$$S = P(\pi/2) = \begin{bmatrix} 1 & 0 \\ 0 & i \end{bmatrix}, \quad S|1\rangle = i|1\rangle, \quad S^2 = Z.$$

S is a quarter-turn around the Bloch sphere's z-axis. It changes $|+\rangle$ into $(|0\rangle + i|1\rangle)/\sqrt{2}$ while leaving basis measurement probabilities unchanged. Its inverse is $S^\dagger = P(-\pi/2)$, not S itself.

```
S -I Q -0 Q
```

Input	Measured bit	Quantum output
0	0	$ 0\rangle$
1	1	$i 1\rangle$

6.6 T phase (T)




$$T = P(\pi/4) = \begin{bmatrix} 1 & 0 \\ 0 & e^{i\pi/4} \end{bmatrix}, \quad T^2 = S, \quad T^4 = Z.$$

T is an eighth-turn around z. Together with gates such as H and CNOT, it is an important non-Clifford resource for expressing general quantum computations. Its inverse is $T^\dagger = P(-\pi/4)$. **Tdg** and **Tinv** lower to that gate, and **Sdg** and **Sinv** lower to $S^\dagger = P(-\pi/2)$. **PHASEdg=pi/4** is the same negative phase.

T -I Q -0 Q

Input	Measured bit	Quantum output
0	0	$ 0\rangle$
1	1	$e^{i\pi/4} 1\rangle$

6.7 Arbitrary phase (PHASE)



$$P(\theta) = \begin{bmatrix} 1 & 0 \\ 0 & e^{i\theta} \end{bmatrix}, \quad \alpha|0\rangle + \beta|1\rangle \mapsto \alpha|0\rangle + \beta e^{i\theta}|1\rangle.$$

PHASE is the general form behind S, T, and Z. It preserves $|\alpha|^2$ and $|\beta|^2$ because $|e^{i\theta}| = 1$, yet rotates the relative phase by θ . On the Bloch sphere this is a z-axis rotation up to global phase. A later H or another basis-changing gate is normally needed to turn that phase difference into changed computational-basis probabilities. The angle is embedded in the opcode:

```
PHASE=1.5708 -I Q -0 Q      # radians
PHASE=3/2 -I Q -0 Q        # rational radians
PHASE=90d -I Q -0 Q        # degrees
PHASE=45/2d -I Q -0 Q      # rational degrees
PHASE=pi -I Q -0 Q
PHASE=-11pi/6 -I Q -0 Q
PHASE=2*pi/3 -I Q -0 Q
PHASEdg=pi/4 -I Q -0 Q      # P(-pi/4)
invPHASE=pi/2 -I Q -0 Q     # P(-pi/2)
```

Accepted numeric atoms include signed integers, decimals, leading-dot decimals, and one optional rational denominator. Pi notation accepts an optional sign, an integer coefficient, an optional *, and an optional positive integer denominator. A zero denominator and malformed angle raise an error.

Input	Measured bit	Quantum output
0	0	$ 0\rangle$
1	1	$e^{i\theta} 1\rangle$

6.8 Dagger and inverse syntax

A primitive gate accepts **dg** (dagger) or **inv** (inverse), written after the opcode or before it. Both spellings set **reverse=true**.

```

Xdg -I Q -0 Q
invH -I Q -0 Q
CNOTdg -I A -0 T
dgSWAP -I A B -0 A B
Sdg -I Q -0 Q          # P(-pi/2)
Tinv -I Q -0 Q          # P(-pi/4)
PHASEdg=pi/4 -I Q -0 Q  # P(-pi/4)

```


Sdg, Sinv, dgS, and invS synthesize $S^\dagger$. Tdg, Tinv, dgT, and invT synthesize $T^\dagger$. PHASEdg=angle, PHASEinv=angle, dgPHASE=angle, and invPHASE=angle negate the angle. For these self-inverse gates the marker emits the ordinary gate, which is already its own inverse:

X, Y, Z, H, CNOT, CCNOT, CZ, CY, SWAP.

A dg or inv marker on a derived gate such as ANDdg is an unknown opcode. MEASUREdg parses and warns, because measurement has no inverse. Compiler-internal RESET remains unavailable in source.

7 Controlled and two-qubit primitive gates

7.1 Controlled NOT (CNOT)



$$|c, t\rangle \mapsto |c, t \oplus c\rangle.$$


CNOT leaves the control unchanged. When the control is zero, the target is unchanged; when the control is one, X acts on the target. Because the map is a permutation of basis states, it is reversible and self-inverse. A superposition control can entangle the wires, so “classical XOR” describes only its basis-state action.

```
CNOT -I Control -0 Target
```

c	t	c'	t'
0	0	0	0
0	1	0	1
1	0	1	1
1	1	1	0

CNOT is a reversible embedding of XOR because it preserves the control. In the workbench, the control comes from **Control A** and the target from **Target particle**; **Control B** is ignored.

7.2 Doubly controlled NOT / Toffoli (CCNOT)



$$|a, b, t\rangle \mapsto |a, b, t \oplus (a \wedge b)\rangle.$$

CCNOT, also called Toffoli, applies X only when both controls are one. It is a reversible universal building block for classical computation embedded in a quantum circuit. The target’s previous value matters: zero computes AND, while one computes NAND. The gate is self-inverse.

CCNOT -I A B -0 Target


a	b	t	a'	b'	t'
0	0	0	0	0	0
0	0	1	0	0	1
0	1	0	0	1	0
0	1	1	0	1	1
1	0	0	1	0	0
1	0	1	1	0	1
1	1	0	1	1	1
1	1	1	1	1	0

With a target initialized to zero, Toffoli computes classical AND reversibly.

Controls are preserved, the target is the output

In the table, the a' and b' columns always equal a and b : the two controls are only read. Only t can change, and it flips exactly in the two rows where $a = b = 1$. In the **Interactive workbench**, choose **CCNOT** in **Gate**, the two inputs in **Control A** and **Control B**, and the output wire in **Target particle**. The *CCNOT demo* starter circuit is exactly this gate with both controls prepared as 1.

7.3 Controlled Z (CZ)



$$\text{CZ}|c, t\rangle = (-1)^{ct}|c, t\rangle, \quad \text{CZ} = \text{diag}(1, 1, 1, -1).$$


CZ adds a minus sign only to the $|11\rangle$ amplitude. No computational-basis bit changes, but superpositions can acquire observable relative phase. CZ is symmetric between its two wires, even though QPU syntax writes one as the input control and one as the output target. It is self-inverse.

CZ -I Control -0 Target

Input	Measured output	Amplitude multiplier	State
00	00	+1	$ 00\rangle$
01	01	+1	$ 01\rangle$
10	10	+1	$ 10\rangle$
11	11	-1	$- 11\rangle$

The classical columns are unchanged; the quantum phase is the operation.

7.4 Controlled Y (CY)




$$|0, t\rangle \mapsto |0, t\rangle, \quad |1, 0\rangle \mapsto i|1, 1\rangle, \quad |1, 1\rangle \mapsto -i|1, 0\rangle.$$

CY applies the Pauli Y rotation to the target only when the control is one. Its basis-bit pattern resembles CNOT, but the $\pm i$ factors distinguish their behavior on superpositions. Since $Y^2 = I$, CY is self-inverse.

CY -I Control -0 Target

c	t	c'	t'	Quantum phase
0	0	0	0	+1
0	1	0	1	+1
1	0	1	1	i
1	1	1	0	$-i$

7.5 Swap (SWAP)



$$\text{SWAP}|a, b\rangle = |b, a\rangle.$$

SWAP exchanges the complete quantum states of two wires, including amplitudes and any correlations with the rest of the register. It does not copy either state. Three CNOTs can decompose SWAP, but the QPU system provides it as one primitive instruction. Applying SWAP twice restores the original ordering.

SWAP -I A B -0 A B

The two output tokens are mandatory and must resolve to the same two wires, in the same order, as the inputs.

a	b	a'	b'
0	0	0	0
0	1	1	0
1	0	0	1
1	1	1	1

8 Derived logical gates

Derived gates are QPU instructions implemented as reversible target mutations. Inputs are preserved. If target t starts in zero, AND/OR/XOR write the familiar classical function. If it does not, the true behavior is the XOR embedding shown below.

One shape for every two-input gate

AND, NAND, OR, and XOR all follow $t' = t \oplus f(a, b)$. The two inputs a and b (**Control A** and **Control B** in the workbench) are read and left unchanged. The target t (**Target particle**) is the output workspace and is the only wire that changes. Only f differs between the four gates. To get the ordinary classical answer, start the target at 0 with SET Out Op.

8.1 Logical NOT (NOT)



$$t' = t \oplus 1.$$


NOT is the derived logical spelling of a bit flip. In the current simulator it uses the same X matrix on the output target. The required input establishes the one-input instruction shape, but a distinct input is not copied into the output; use the same token on both sides.

NOT -I Target -O Target

t	t'
0	1
1	0

The current gate implementation mutates the output target. Use the same token for input and output; a distinct input token is not copied.

8.2 Reversible AND (AND)




$$t' = t \oplus (a \wedge b).$$

AND flips the target exactly on the basis components where both controls are one. This is Toffoli-style behavior. It remains linear on superpositions and may entangle the target with its controls. Calling it AND is convenient only when the target is deliberately initialized.

SET Out Op
AND -I A B -O Out

a	b	$t = 0$	$t' = a \wedge b$
0	0	0	0
0	1	0	0
1	0	0	0
1	1	0	1

8.3 Reversible NAND (NAND)




$$t' = t \oplus (a \wedge b) \oplus 1.$$

NAND first performs the all-controls-active mutation and then inverts the target. With target zero it produces the familiar NAND column. With target one it instead produces AND, illustrating why reversible target state must always be included in the full mathematical definition.

SET Out Op
NAND -I A B -O Out

a	b	$t = 0$	$t' = \neg(a \wedge b)$
0	0	0	1
0	1	0	1
1	0	0	1
1	1	0	0

8.4 Reversible OR (OR)



$$t' = t \oplus (a \vee b).$$


OR uses an “any control is one” predicate. If either or both controls are active, the target is flipped. This direct derived operation is equivalent on basis states to combining XOR and AND through $a \vee b = a \oplus b \oplus (a \wedge b)$.

```
SET Out Op
OR -I A B -O Out
```

a	b	$t = 0$	$t' = a \vee b$
0	0	0	0
0	1	0	1
1	0	0	1
1	1	0	1

For bits, $a \vee b = a \oplus b \oplus (a \wedge b)$.

8.5 Reversible XOR (XOR)



$$t' = t \oplus a \oplus b.$$


XOR flips the target when an odd number of controls is one. With two inputs, this is the ordinary exclusive-or predicate. The parity definition explains why the 11 row does not flip: two active controls have even parity.

```
SET Out Op
XOR -I A B -O Out
```

a	b	$t = 0$	$t' = a \oplus b$
0	0	0	0
0	1	0	1
1	0	0	1
1	1	0	0

9 Measurement gate

9.1 Measurement (MEASURE)



```
MEASURE -I Q      # exactly one named wire
MEASURE           # every wire known at this point
```

MEASURE takes no -O. Repeated one-wire commands are useful when only declared returns should become classical. A bare command emits one measurement operation per currently allocated wire.

State	Outcome probabilities	Collapsed state
$ 0\rangle$	$P(0) = 1$	$ 0\rangle$
$ 1\rangle$	$P(1) = 1$	$ 1\rangle$
$\alpha 0\rangle + \beta 1\rangle$	$P(0) = \alpha ^2, P(1) = \beta ^2$	$ 0\rangle$ or $ 1\rangle$

Three ways to measure in the builder

The meter A step *in the circuit*, like any other gate. It runs every time the circuit runs and is saved in the protocol as MEASURE.

Measure target Measures the selected **Target particle** once, right now. Nothing is added to the circuit.

Measure all Measures every wire not yet measured, right now.

MEASURE, like RESET, is not reversible: once a superposition collapses, the lost amplitudes cannot be recovered. Choose **Reset state** to rewind and sample again.

10 Reset: an operation that is not a gate

The word *reset* appears in two unrelated places: the compiler's internal RESET operation, and three buttons in the Circuit builder. Neither is an ordinary quantum gate.

10.1 Reversible gates versus RESET

Every gate in the palette except the meter is *reversible*: its effect can always be undone. X, Y, Z, H, CNOT, CCNOT, CZ, CY, SWAP, NOT, AND, NAND, OR, and XOR undo themselves when applied a second time. S, T, and PHASE are undone by the same phase with the opposite sign. Mathematically each is a *unitary* matrix U , and $U^\dagger$ undoes it.

RESET is different. It forces one wire to $|0\rangle$ whatever it held before:

$$|0\rangle \mapsto |0\rangle, \quad |1\rangle \mapsto |0\rangle, \quad \alpha|0\rangle + \beta|1\rangle \mapsto |0\rangle.$$

Two different inputs, $|0\rangle$ and $|1\rangle$, give the same output, so nothing can tell afterwards which one was there. The information has been erased, and no matrix U can undo that. RESET is therefore not unitary, just like MEASURE.

	Reversible gates (X, H, MEASURE, CNOT, AND, ...)		RESET
Unitary?	Yes	No	No
Can be undone?	Yes, by the inverse gate	No	No
Result depends on chance?	No	Yes	No
Written in source?	Yes	Yes	No — write SET wire 0p
Shown on canvas?	Yes	Yes, as the meter	No (hidden)

10.2 Where RESET comes from

You never type RESET. The compiler inserts it when a wire must start at zero:

- SET Out 0p on a workspace wire. (SET Out 1p and SET Out sp are instead prepared by applying X or H to a wire that starts at 0.)
- INCREASECYCLE, which zeroes pending workspace before the next stage.
- RUNCHILD and CALL, which zero each parent output wire once before expanding the child.

RESET steps are hidden from the canvas and filtered from the **Execution log** so that the visible circuit shows only the gates you wrote. They still run during simulation.

Reset a wire only when it is not entangled

In this simulator, RESET keeps the part of the state in which the wire is already 0 and rescales it; only if that part is empty does it move the $|1\rangle$ amplitudes to $|0\rangle$. On a fresh wire this is exactly “prepare $|0\rangle$ ”. On a wire entangled with others it also changes the other wires. For example, resetting one half of a Bell pair forces the other half to 0 as well. In multi-stage circuits, *uncompute* temporary wires back to 0 with reversible gates instead of relying on a reset.

10.3 Reset buttons compared

None of these buttons inserts a RESET operation into your circuit. They only change what the page is showing and simulating.

Kept or cleared?	Reset state	Clear circuit	Reset site	Notes
Gates on the canvas	kept	cleared	cleared	
Number of wires	kept	kept	back to 3	
Start states (<i>name</i> start)	kept	kept	all 0p	
Protocol editor text	kept	rewritten	default	Clear circuit writes an empty canvas protocol into the editor.
Compiled PARAMS / RETURNVALS	kept	cleared	cleared	
Quantum state and measurements	rewound	rewound	rewound	Back to the start states; log restarts.
Workbench selections	kept	kept	defaults	Gate H, target q0, controls q1 and q2, phase 90°.
Custom gates and catalog	kept	kept	kept	Nothing removes these except Remove or closing the browser session.

Which one do I want?

Reset state “Run the same experiment again.” Use it between repeated runs of a circuit that measures, to collect several samples.

Clear circuit “Keep my wires, start a new circuit.”

Reset site “Put everything back the way it was when I first opened the page.” The **Reset site** entry in the navigation menu and **Reset site completely** on the **More** page do the same thing.

11 Glossary

Term	Meaning
Amplitude	A complex coefficient of a basis state. Its squared magnitude contributes to measurement probability.

Term	Meaning
Ancilla	An auxiliary qubit or workspace wire, usually prepared in a known state, used to hold intermediate information in a reversible computation.
Basis state	A state used as a coordinate direction. The computational basis for one qubit is $ 0\rangle, 1\rangle$.
Bloch sphere	A geometric representation of pure one-qubit states after ignoring global phase. North/south are $ 0\rangle/ 1\rangle$; longitude represents relative phase.
Boolean algebra	The mathematics of 0/1 values and operations such as NOT, AND, OR, and XOR.
Child process	A separately named protocol expanded into a parent through RUNCHILD or CALL.
Combinational logic	A classical circuit whose output depends only on current inputs, not stored history.
Control	A wire whose basis condition decides whether a controlled operation acts on its target. The control itself is preserved by the standard controlled gates.
Entanglement	A joint multi-qubit state that cannot be written as a tensor product of independent qubit states.
Fan-out	Sending one classical value to multiple destinations. Unknown quantum states cannot be universally copied; CNOT copies only known basis information onto a zero target and otherwise may entangle.
Gate	A state transformation. Closed-system quantum gates are unitary; measurement and reset have different semantics.
Garbage	Reversible-workspace information no longer wanted as output but still present because it cannot simply be erased.
Global phase	A unit complex multiplier applied to the whole state. It does not change physical measurement predictions.
Ket	Dirac notation $ \psi\rangle$ for a state column vector.
Measurement	Conversion of quantum amplitudes into a sampled classical result, with state collapse and renormalization.
Phase	The angular part of a complex amplitude. Relative phase controls interference even when immediate basis probabilities are unchanged.
Process	A named QPU protocol body with ordered PARAMS and RETURNVALS.
Protocol	Plain-text QPU instructions describing preparation, gates, composition, measurement, and returned references.
Qubit	A normalized two-component complex state $\alpha 0\rangle + \beta 1\rangle$, not merely an unknown classical bit.
Register	One or more logical wires considered together. Current JOIN/SPLIT syntax is accepted but does not yet lower to register operations.
Relative phase	The phase difference between amplitudes. Unlike global phase, it can change later interference and measurement probabilities.
Superposition	A coherent linear combination of basis states. <code>sp</code> prepares the equal plus superposition.
Target	The wire whose amplitudes a gate mutates; normally the first <code>-0</code> reference.
Tensor product	The operation $\otimes$ that combines state spaces. Two one-qubit vectors produce a four-amplitude two-qubit vector.
Token	A source-level logical name resolved to a qubit wire or alias during compilation.
Truth table	Listed input/output expectations for measured RETURNVALS. It describes basis behavior but not arbitrary phase or entanglement.
Uncomputation	Running suitable inverse gates in reverse order to return temporary workspace to a known state without irreversible erasure.
Unitary	A matrix satisfying $U^\dagger U = I$, preserving normalization and having a reversible inverse.